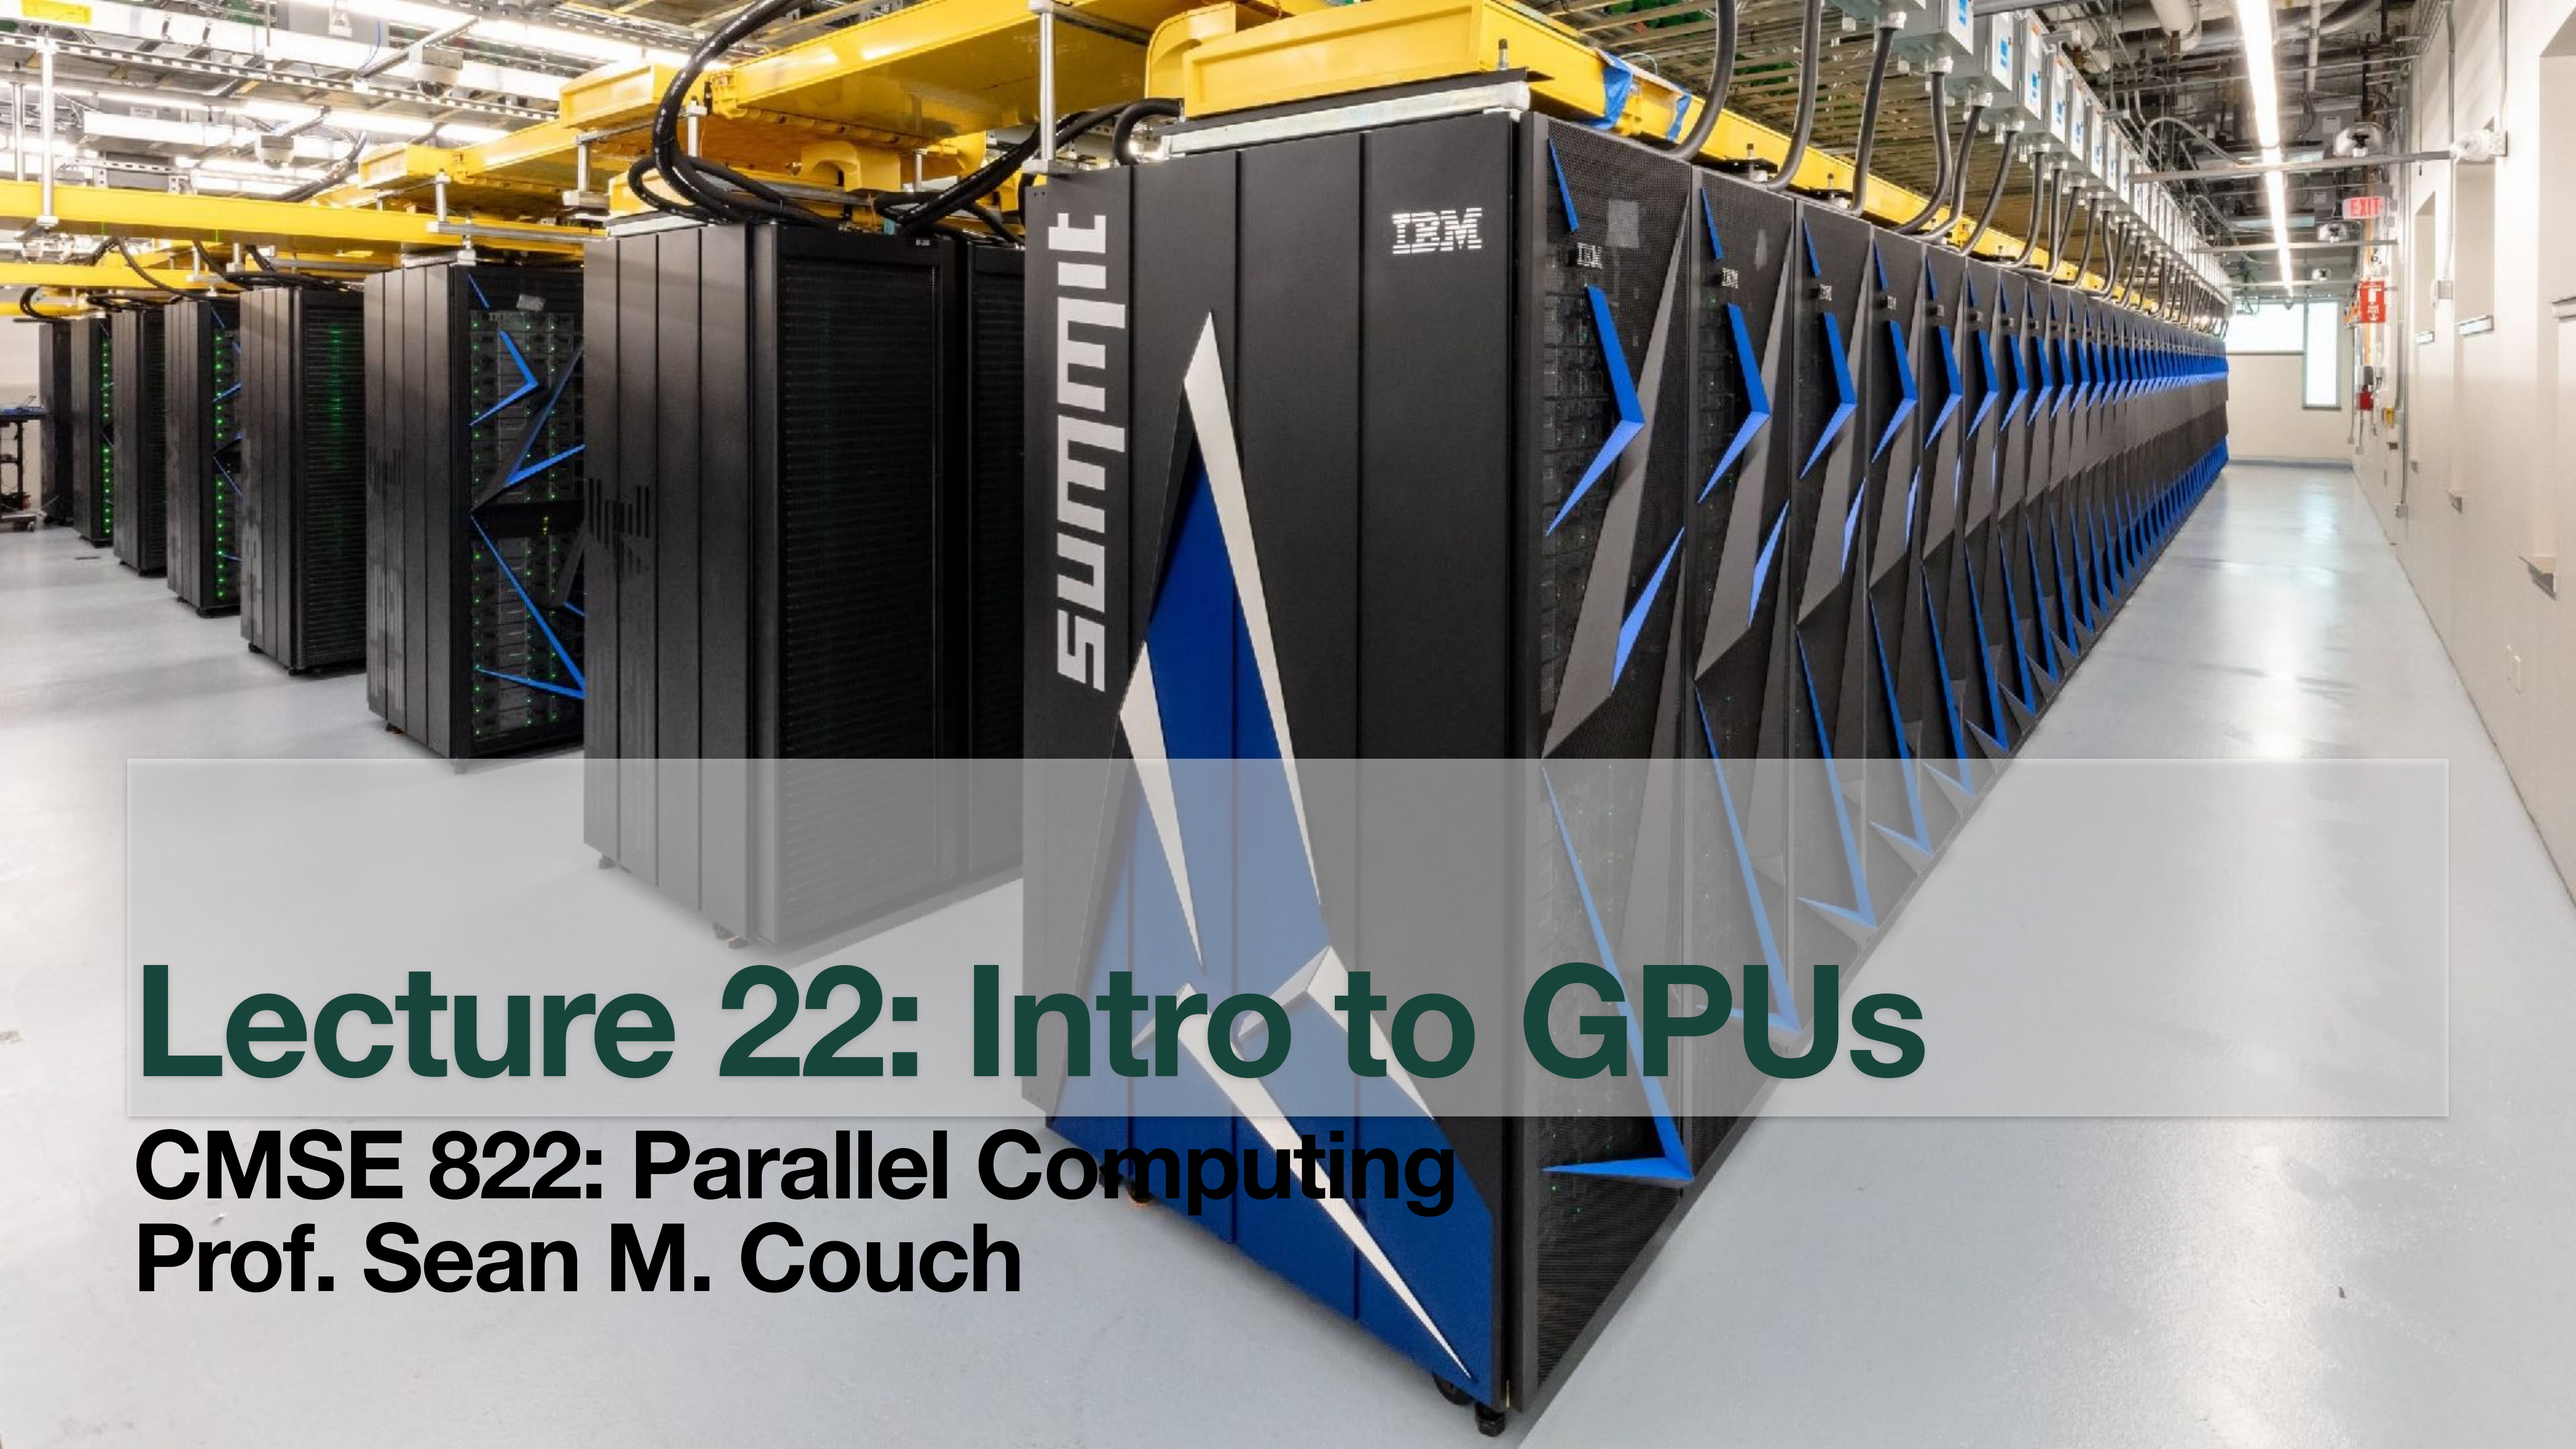

Lecture 22: Intro to GPUs

CMSE 822: Parallel Computing

Prof. Sean M. Couch

Today's Plan

- Fill out the Project Spreadsheet here (see #general channel).
- Questions
- Group work on projects
- Break!

Puppy time!

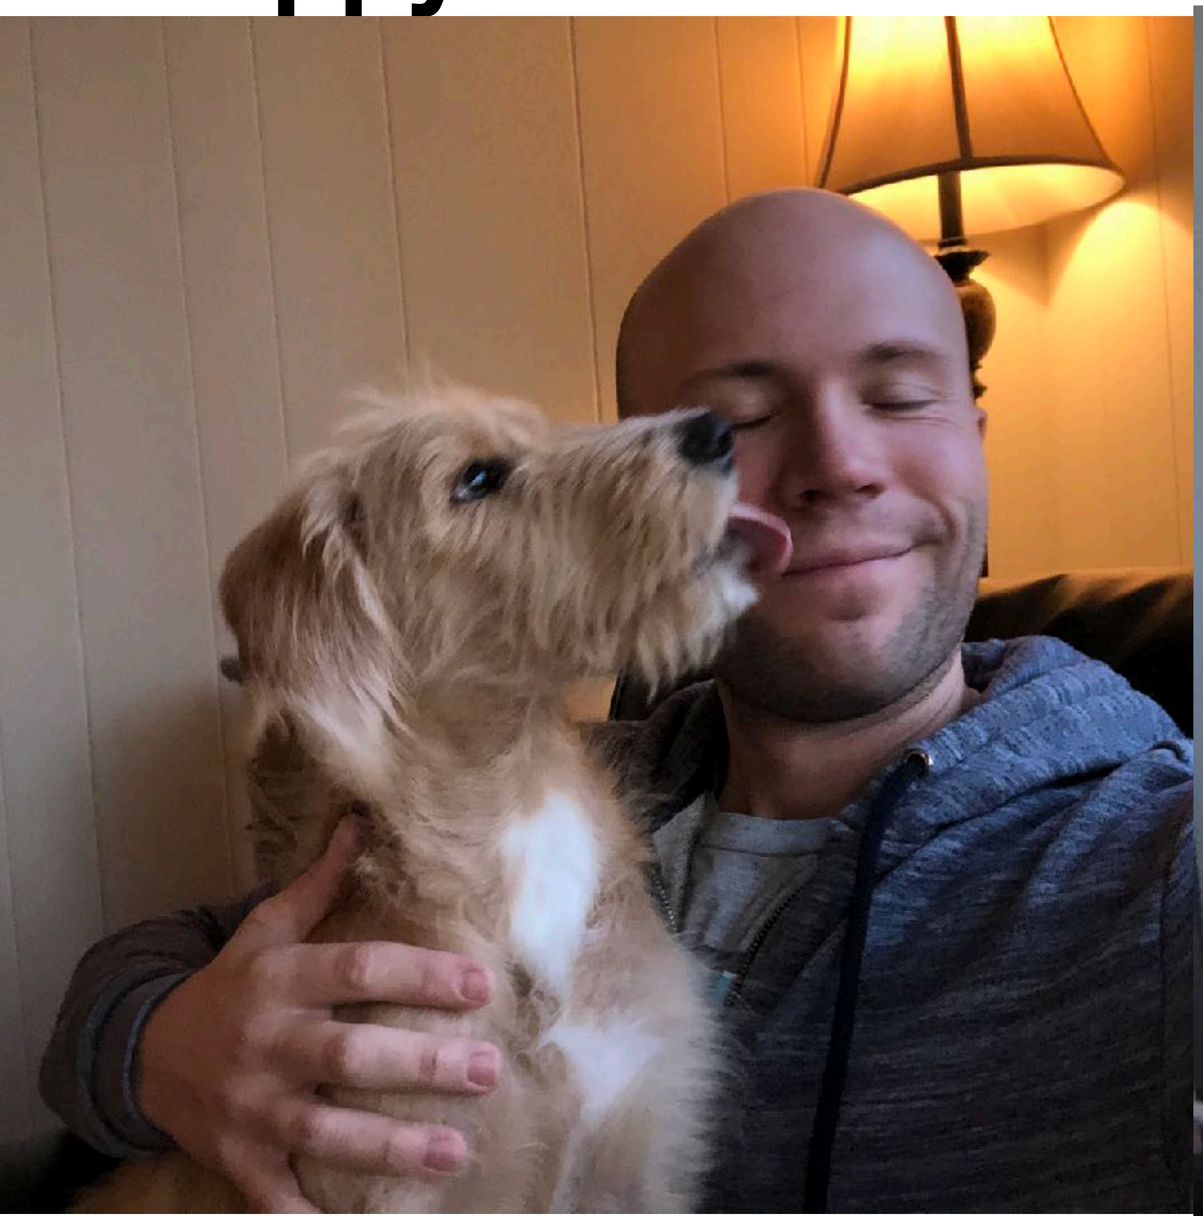

PCA Questions

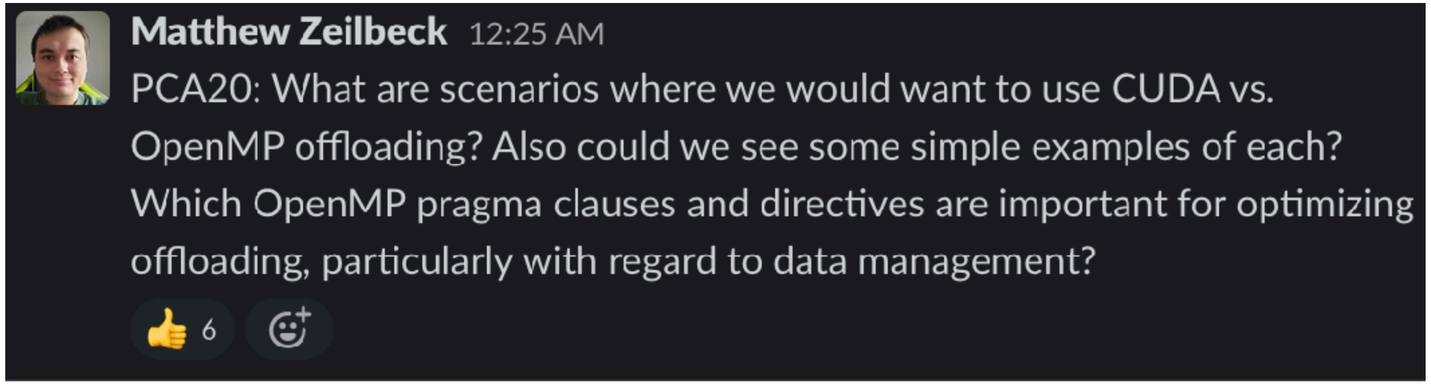

- In general, CUDA will give better performance but at the cost of portability

PCA Questions

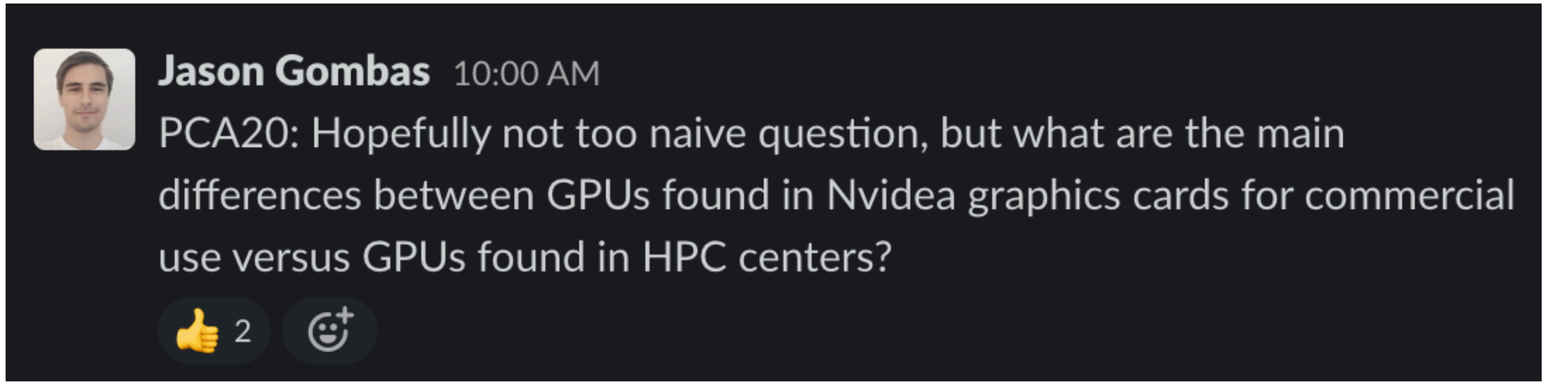

- GPGPUs optimized for floating point (single AND double precision)
- GPGPUs have much better memory bandwidth (big deal for scientific computing)

PCA Questions

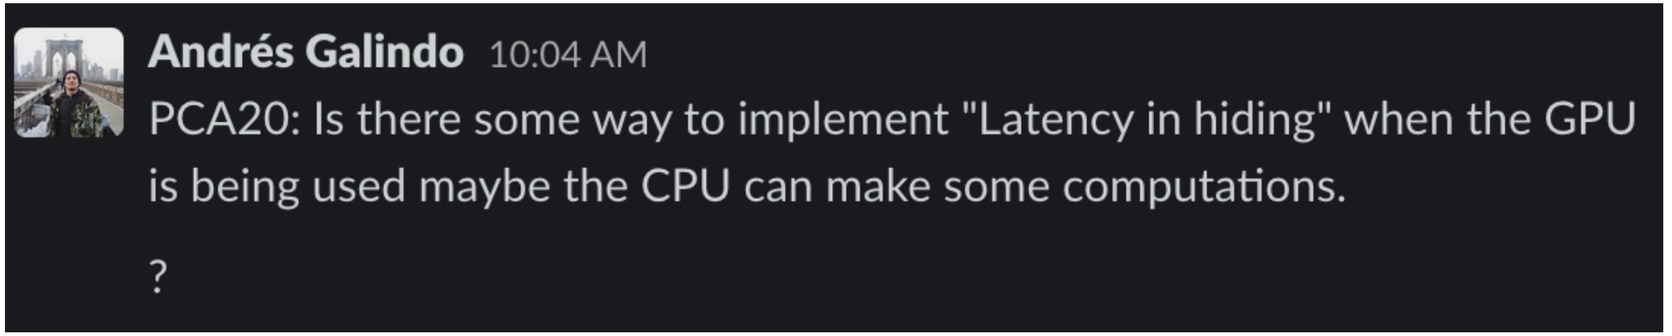

- Yes! Critical in some applications.
- Task-based approach helpful

Project work!